

# OSS security, here be dragons!

Fermin J. Serna - @fjserna



**Chief Security Officer @ Semmle**  
**Distinguished Engineer @ GitHub**

**Previously:**

- Head of Product Security @ Google

**Other:**

- Two times Pwnie Award nomination
  - 2016 winner: glibc getaddrinfo()
- EMET initial main developer
- I once owned Charlie Miller's iPhone at Pwn2Own



**@fjserna**



## The Problem


OSS has won! What about security?



# 99% of software uses open source

**Open Source has won!**  
**But it comes with challenges:**

- Code quality
- Supply chain
- Common development best practices:
  - Code review
  - Testing
  - ...
- Security: CVE tagging



# CVEs in OSS



## **Common Vulnerabilities and Exposures**

*The Standard for Information Security Vulnerability Names*

A test of a random ImageMagick vulnerability against Ubuntu Xenial shows that it, indeed, continues to reproduce.

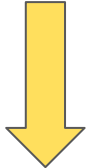
This is in addition to the >100 security bugs OSS-Fuzz found and publicly disclosed due to hitting their disclosure deadline, and which still have not been fixed [3].

<https://www.openwall.com/lists/oss-security/2019/06/15/2>



# Linux Kernel Backdoor (2003)

```
@@ -1111,6 +1111,8 @@
        schedule();
        goto repeat;
    }
+   if ((options == (__WCLONE|__WALL)) && (current->uid = 0))
+       retval = -EINVAL;
    retval = -ECHILD;
end_wait4:
    current->state = TASK_RUNNING;
```




# Supply chain problems...

## **NPM Worm Vulnerability Disclosed**

The NPM project has [formally acknowledged a long-standing security vulnerability](#) in which it is possible for malicious packages to run arbitrary code on developer's systems, leading to the first NPM created worm. In [vulnerability note VU319816](#), titled

<https://www.infoq.com/news/2016/03/npm-infection/>



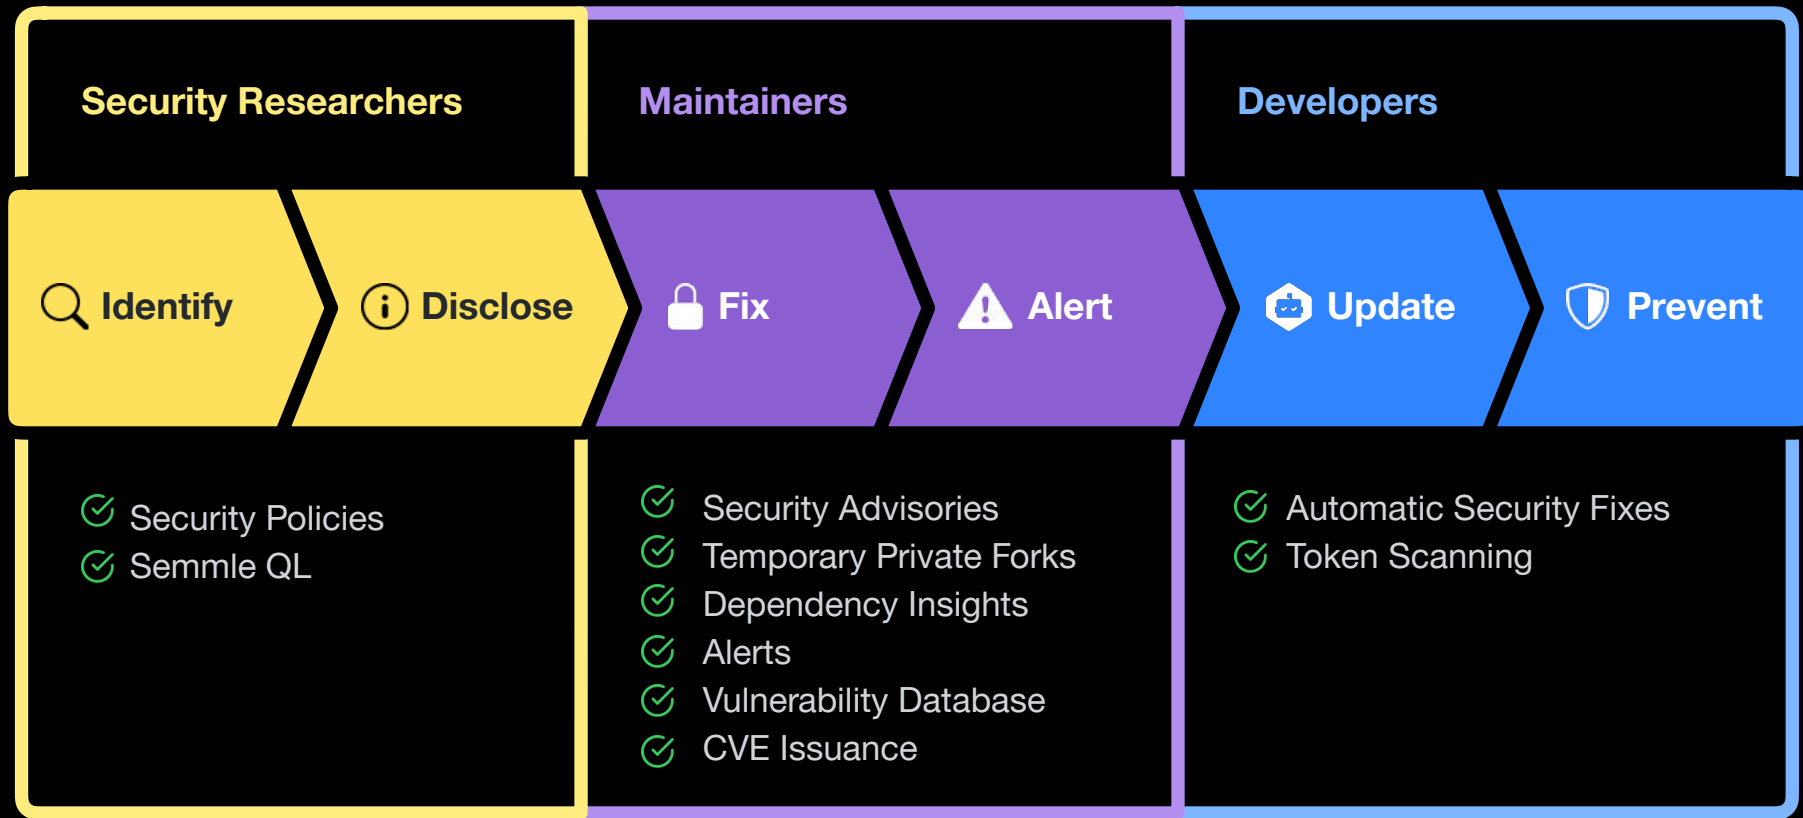




What is Github doing?



# Open Source Security Workflow



# Semmler QL - Automating Variant Analysis

QL is an object oriented **query language** for exploring code as data

## Treat code as data

Security engineers use QL to

- Explore code iteratively
- Map attack surfaces
- Automate variant analysis

```
from AddExpr a, Variable v, RelationalOperation cmp
where a.getAnOperand() = v.getAnAccess()
      and cmp.getAnOperand() = a
      and cmp.getAnOperand() = v.getAnAccess()
      and forall(Expr op | op = a.getAnOperand() |
                 op.getType().getSize() < 4)
      and not a.getExplicitlyConverted().getType().getSize() < 4
select cmp, "Bad overflow check"
```

*Sample QL query identifying common mistakes made checking for integer overflows*



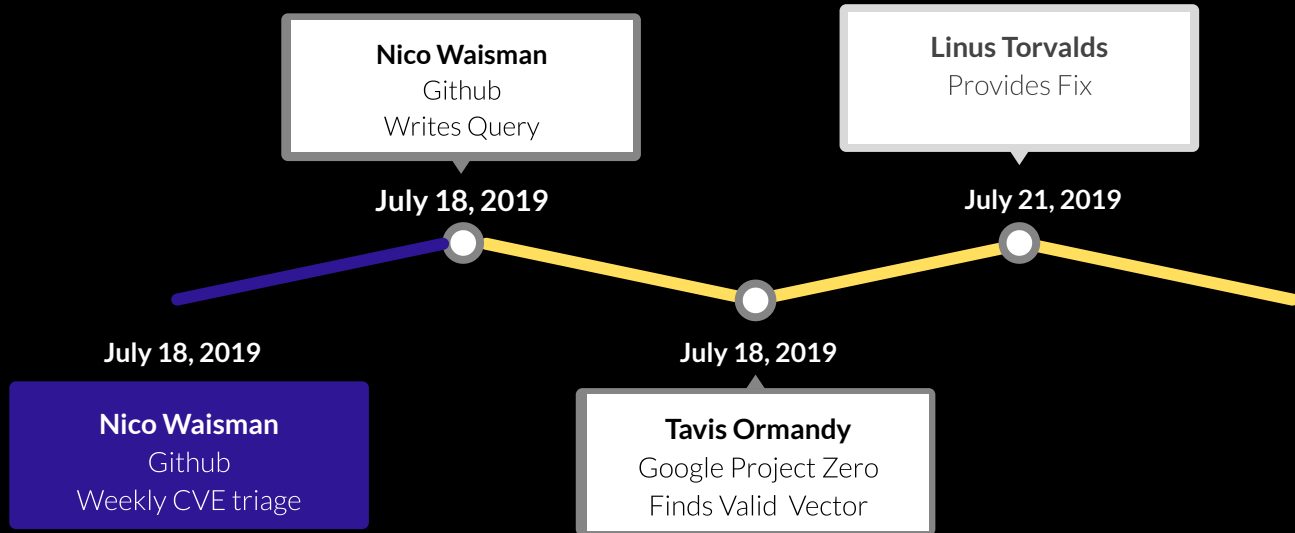
What is Github “also” doing?

Variant Analysis (CVE triage)



# Linux Kernel

## Buffer Overflow



# Array write inside loop

## We are looking for:

```
while (condition) {  
    [...]  
  
    array[x] = y;  
  
    [...]  
}
```

```
import cpp  
  
from Variable buffer, Expr bufferWriteBase, ArrayExpr bufferWrite,  
Assignment a, Loop loop  
where  
    buffer.getType() instanceof ArrayType and  
    bufferWriteBase = buffer.getAnAccess() and  
    bufferWriteBase = bufferWrite.getArrayBase() and  
    bufferWrite = a.getLValue() and  
    loop = a.getEnclosingStmt().getParent*()  
select a
```





Nico Waisman

@nicowaisman

Follow

Now you can go and hunt bugs in Linux kernel here:

[lgtm.com/query/46597517...](https://lgtm.com/query/46597517...)

Warning: There will be a LOT of results.

12:03 PM - 18 Jul 2019



Tavis Ormandy  @taviso · Jul 18

Replying to @nicowaisman

That is kinda fun, there are some real ones in there, although I dunno how many users plug in malicious monitors 😄 (I think the length of the last pair can exceed the edid[2] length check, i.e. a ton of len=1 and a len=0x1f at the end)

[lgtm.com/projects/g/tor...](https://lgtm.com/projects/g/tor...)



5



4



18



```
Date: Sun, 21 Jul 2019 11:03:01 -0700
From: Linus Torvalds <torvalds@...ux-foundation.org>
To: Tavis Ormandy <taviso@...il.com>, Bartlomiej Zolnierkiewicz <b.zolnierkie@...
    Daniel Vetter <daniel.vetter@...ll.ch>
Cc: oss-security@...ts.openwall.com
Subject: Re: stack buffer overflow in fbdev
```

Completely untested patch attached. There are probably better ways to do this.

Adding the proper people to the cc, and quoting Tavis' email in its entirety.

Daniel - you got added despite not being explicitly listed as maintainer because you've touched fbdev/core/ more than most lately, plus you know edid anyway. As such: "tag, you're it, sucker".

Linus



# Linux Kernel, never plug untrusted monitors...

```
while (pos < edid[2]) {
    u8 len = edid[pos] & 0x1f, type = (edid[pos] >> 5) & 7;
    pr_debug("Data block %u of %u bytes\n", type, len);
    if (type == 2) {
        for (i = pos; i < pos + len; i++) {
            u8 idx = edid[pos + i] & 0x7f;
            svd[svd_n++] = idx;
            pr_debug("N%sative mode %#d\n",
                    edid[pos + i] & 0x80 ? "" : "on-", idx);
        }
    } else if (type == 3 && len >= 3) {
        /* Check Vendor Specific Data Block. For HDMI,
           it is always 00-0C-03 for HDMI Licensing, LLC. */
        if (edid[pos + 1] == 3 && edid[pos + 2] == 0xc &&
            edid[pos + 3] == 0)
            specs->misc |= FB_MISC_HDMI;
    }
    pos += len + 1;
}
```






# U-Boot: A variant analysis journey



# What is U-Boot?

- Open source bootloader
- Used by
  - Kindle devices
  - ChromeOS ARM devices
  - IoT
- Supports verifiable boot
  - Any vulnerability before the signature check is a potential jailbreak
  - Filesystems should then be considered untrusted
  - .....Not to mention any networking :)



### Seed Vulnerability:

- 2x memcpy() size argument is attacker controlled
- size comes from the NFS packet
- No size constraints

```
static int nfs_readlink_reply(uchar *pkt, unsigned len)
{
    [...]

    /* new path length */
    rlen = ntohl(rpc_pkt.u.reply.data[1 + nfsv3_data_offset]);

    if (((char *)&(rpc_pkt.u.reply.data[2 + nfsv3_data_offset])) != '/') {
        int pathlen;

        strcat(nfs_path, "/");
        pathlen = strlen(nfs_path);
        memcpy(nfs_path + pathlen,
               (uchar *)&(rpc_pkt.u.reply.data[2 + nfsv3_data_offset]),
               rlen);
        nfs_path[pathlen + rlen] = 0;
    } else {
        memcpy(nfs_path,
               (uchar *)&(rpc_pkt.u.reply.data[2 + nfsv3_data_offset]),
               rlen);
        nfs_path[rlen] = 0;
    }

    return 0;
}
```



## Query initial steps: find all memcpy callers

- **596** instances. Find usage of `rpc_pkt.u.reply`
- **191** instances

```
import cpp

from FunctionCall call
where call.getTarget().getName() = "memcpy"
select call
```

```
import cpp

from FieldAccess access, Field f
where f = access.getTarget() and f.hasName("reply")
select access
```



Given a (problem-specific) set of sources and sinks, is there a path in the **data flow graph** from some source to some sink?



### Query refining process

- Data Flow analysis
- Source: nfs data packet read from the socket
- Sink: memcpy()

### Results:

- 4 instances
- We found the 2 original seed vulnerabilities

```
import cpp
import semmle.code.cpp.dataflow.TaintTracking

class NetworkToMemFuncLength extends TaintTracking::Configuration {
  NetworkToMemFuncLength() { this = "NetworkToMemFuncLength" }

  override predicate isSource(DataFlow::Node source) {
    exists (FieldAccess access, Field f |
      source.asExpr() = access and f = access.getTarget() and
      f.hasName("reply"))
  }

  override predicate isSink(DataFlow::Node sink) {
    exists (FunctionCall fc, Expr argument |
      sink.asExpr() = argument and (argument = fc.getArgument(2)
      and fc.getTarget().hasQualifiedName("memcpy")))
  }
}

from Expr socket_buffer, Expr sizeArg, NetworkToMemFuncLength
  config
where config.hasFlow(DataFlow::exprNode(socket_buffer),
  DataFlow::exprNode(sizeArg))
select sizeArg
```



# Findings triage

```
filefh3_length = ntohl(rpc_pkt.u.reply.data[1]);  
if (filefh3_length > NFS3_FHSIZE)  
    filefh3_length = NFS3_FHSIZE;  
  
memcpy(filefh, rpc_pkt.u.reply.data + 2, filefh3_length);
```

Looks like they are doing the correct by checking the length

Until you discover that filefh3\_length is a signed integer!!!

**Vulnerability Counter += 1**



## Findings triage

Same story:

- rlen comes directly from the NFS packet
- rlen is consumed as the size for the memcpy() operation
- No checks!

memcpy() happens inside store\_block() in two different locations

```

static int nfs_read_reply(ucred *ucred, struct nfsd *nfsd, struct rpc_pkthdr *pkt, unsigned len)
{
    [...]

    if (supported_nfs_versions & NFSV2_FLAG) {
        rlen = ntohl(rpc_pkt.u.reply.data[18]); // <-- rlen is attacker-controlled
        data_ptr = (uchar *)&(rpc_pkt.u.reply.data[19]);
    } else { /* NFSV3_FLAG */
        int nfsv3_data_offset =
            nfs_get_attributes_offset(rpc_pkt.u.reply.data);

        /* count the size */
        rlen = ntohl(rpc_pkt.u.reply.data[1 + nfsv3_data_offset]); // <-- rlen is attacker-controlled
        /* Skip unused values :
           EOF:          32 bits value,
           data_size:   32 bits value,
           ... */
        data_ptr = (uchar *)
            &(rpc_pkt.u.reply.data[4 + nfsv3_data_offset]);
    }

    if (store_block(data_ptr, nfs_offset, rlen)) // <-- We pass to store_block the attacker-controlled rlen
        return -9999;

    [...]
}

```



## Let's think bigger beyond NFS:

- Data Flow analysis
- Source: ntohl(), ntohs() and friends
- Sink: memcpy()

## Results:

- 8 instances
- Finds the all previous vulnerabilities

```
import cpp
import semmle.code.cpp.dataflow.TaintTracking

class NetworkByteOrderTranslation extends Expr {
  NetworkByteOrderTranslation() {
    this = any(MacroInvocation mi |
mi.getOutermostMacroAccess().getMacroName().regexMatch("(?i)(^|.|*_|nt
oh(1|ll|s)")
    ).getExpr()
  }
}

[...]

  override predicate isSource(DataFlow::Node source) {
    exists (FieldAccess access, Field f |
      source.asExpr() instanceof NetworkByteOrderTranslation
    )
  }

[...]


from Expr socket_buffer, Expr sizeArg, NetworkToMemFuncLength config
where config.hasFlow(DataFlow::exprNode(socket_buffer),
DataFlow::exprNode(sizeArg))
select sizeArg
```



# Findings triage

```
#if defined(CONFIG_NETCONSOLE) && !defined(CONFIG_SPL_BUILD)
    nc_input_packet((uchar *)ip + IP_UDP_HDR_SIZE,
                   src_ip,
                   ntohs(ip->udp_dst),
                   ntohs(ip->udp_src),
                   ntohs(ip->udp_len) - UDP_HDR_SIZE); // <- integer
#endif

/*
 * IP header OK. Pass the packet to the current handler.
 */
(*udp_packet_handler)((uchar *)ip + IP_UDP_HDR_SIZE,
                      ntohs(ip->udp_dst),
                      src_ip,
                      ntohs(ip->udp_src),
                      ntohs(ip->udp_len) - UDP_HDR_SIZE); // <- in
```



TCP/IP stack

2 integer underflows with no bounds checking later ending up in memcpy()

**Vulnerability Counter += 2**



# Other vulnerabilities

```
static int nfs_readlink_reply(uchar *pkt, unsigned len)
{
    struct rpc_t rpc_pkt;

    [...]

    memcpy((unsigned char *)&rpc_pkt, pkt, len);
}
```

Plain stack overflow

Also happens in 4 other different functions

**Vulnerability Counter += 5**



## Other vulnerabilities

```
static int nfs_read_reply(uchar *pkt, unsigned len)
{
    struct rpc_t rpc_pkt;

    [...]

    memcpy(&rpc_pkt.u.data[0], pkt, sizeof(rpc_pkt.u.reply));
}
```

Clearly, someone is trying to prevent the overflow

But, forgot to check the lower bound... would be a READ OOB

**Vulnerability Counter += 1**



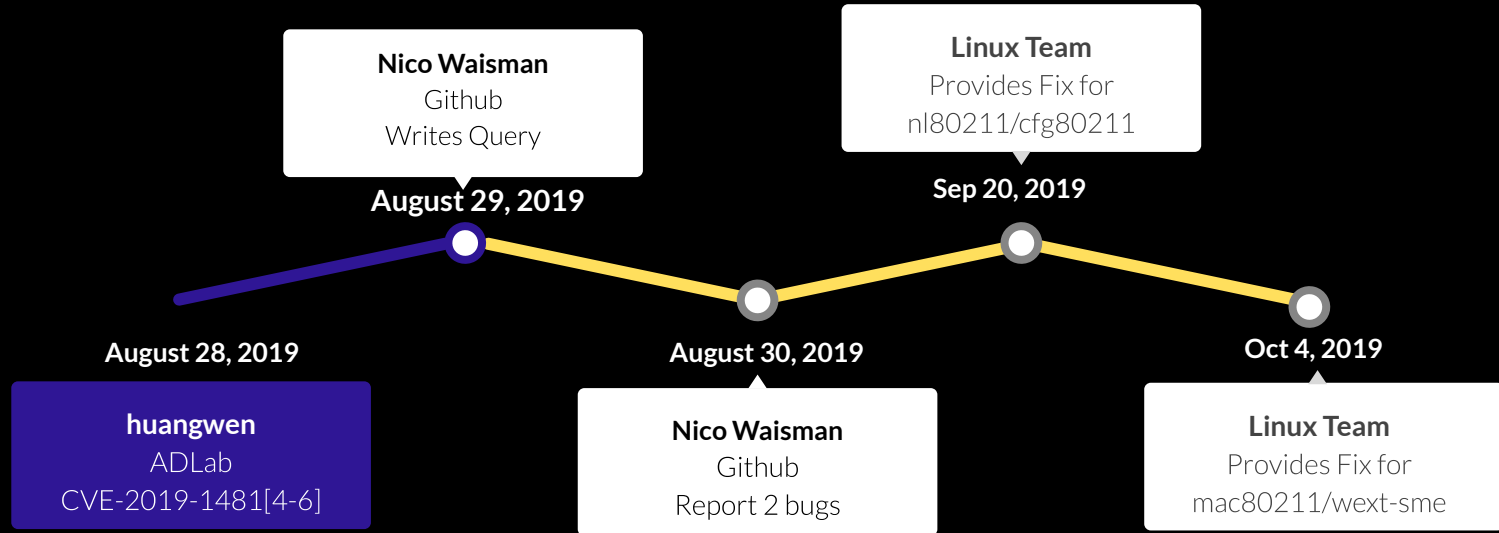
**CVE-2019-14192, CVE-2019-14193, CVE-2019-14194,  
CVE-2019-14195, CVE-2019-14196, CVE-2019-14197,  
CVE-2019-14198, CVE-2019-14199, CVE-2019-14200,  
CVE-2019-14201, CVE-2019-14202, CVE-2019-14203  
and CVE-2019-14204**



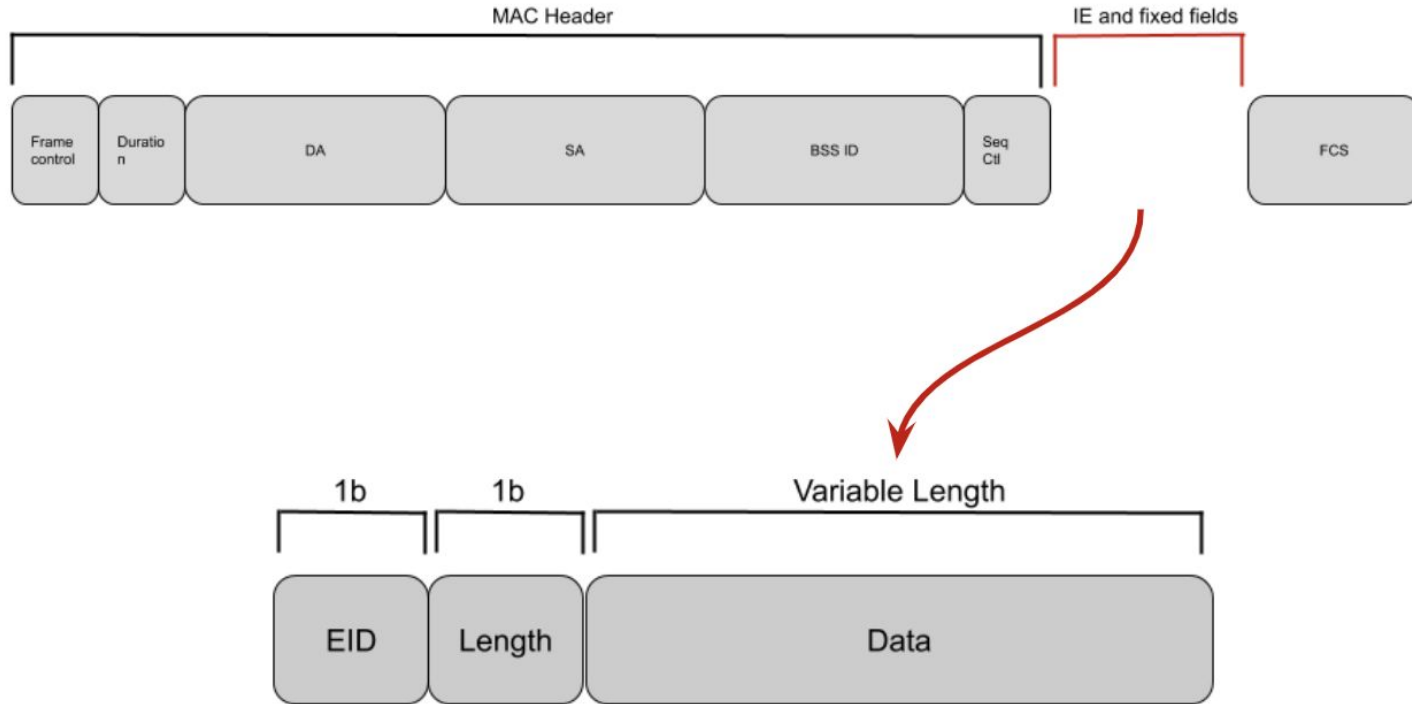
# Wifi drivers and trusting lengths



# Linux Kernel Buffer Overflow



# Crash course on 802.11 framing





## CVE-2019-1481[4-6]

Found by huangwen of  
ADLab of Venustech

```
257 mwifiex_set_uap_rates(struct mwifiex_uap_bss_param *bss_cfg,
258                     struct cfg80211_ap_settings *params)
259 {
260     struct ieee_types_header *rate_ie;
261     int var_offset = offsetof(struct ieee80211_mgmt, u.beacon.variable);
262     const u8 *var_pos = params->beacon.head + var_offset;
263     int len = params->beacon.head_len - var_offset;
264     u8 rate_len = 0;
265
266     rate_ie = (void *)cfg80211_find_ie(WLAN_EID_SUPP_RATES, var_pos, len);
267     if (rate_ie) {
268         memcpy(bss_cfg->rates, rate_ie + 1, rate_ie->len);
269         rate_len = rate_ie->len;
270     }
271
272     rate_ie = (void *)cfg80211_find_ie(WLAN_EID_EXT_SUPP_RATES,
273                                     params->beacon.tail,
274                                     params->beacon.tail_len);
275     if (rate_ie)
276         memcpy(bss_cfg->rates + rate_len, rate_ie + 1, rate_ie->len);
277
278     return;
279 }
```



## We start with cfg80211\_find\_ie:

- Data Flow analysis
- Source: cfg80211\_find\_ie
- Sink: memcpy() size

## Results:

- 13 instances
- No bugs. Code looks buggy, but there were sanitized
- What other function deal with IE?

```
class GetIE extends TaintTracking::Configuration {
  GetIE() { this="cfg80211_find_ie" }

  override predicate isSource(DataFlow::Node source) {
    exists( Function s1 |
      source.asExpr().(FunctionCall).getTarget() = s1
      and
      s1.hasQualifiedName("cfg80211_find_ie")
    )
  }
  override predicate isSink(DataFlow::Node sink) {
    exists(FunctionCall fc |
      sink.asExpr() = fc.getArgument(2)
      and
      fc.getTarget().hasQualifiedName("memcpy")
    )
  }
}
```



## A new IE based function:

- Data Flow analysis
- Source: ieee80211\_bss\_get\_ie
- Sink: memcpy() size

## Results:

- 10 instances
- Most of the match looks promising, but a big chunk where sanitized.

```
class GetIE extends TaintTracking::Configuration {
  GetIE() { this="ieee80211_bss_get_ie" }

  override predicate isSource(DataFlow::Node source) {
    exists( Function s1 |
      source.asExpr().(FunctionCall).getTarget() = s1
      and
      s1.hasQualifiedName("ieee80211_bss_get_ie")
    )
  }

  override predicate isSink(DataFlow::Node sink) {
    exists(FunctionCall fc |
      sink.asExpr() = fc.getArgument(2)
      and
      fc.getTarget().hasQualifiedName("memcpy")
    )
  }
}
```



# CVE-2019-16746

Vulnerability counter: 1

```
1286
1287     if (!conf->ibss_joined) {
1288         const u8 *ssidie;
1289         rcu_read_lock();
1290         ssidie = ieee80211_bss_get_ie(bss, WLAN_EID_SSID);
1291         if (ssidie) {
1292             join.ssid_len = ssidie[1];
1293             memcpy(join.ssid, &ssidie[2], join.ssid_len);
1294         }
1295         rcu_read_unlock();
```

- cw1200 wireless driver
- Remotely exploitable through a beacon



# CVE-2019-17133

Vulnerability counter: 2

```
215
216     rcu_read_lock();
217     ie = ieee80211_bss_get_ie(&wdev->current_bss->pub,
218                             WLAN_EID_SSID);
219     if (ie) {
220         data->flags = 1;
221         data->length = ie[1];
222         memcpy(ssid, ie + 2, data->length);
223     }
```

- cfg80211 wext compat for managed mode.
- Remotely exploitable through a beacon



# Q & A

Thanks!  
@fjserna

